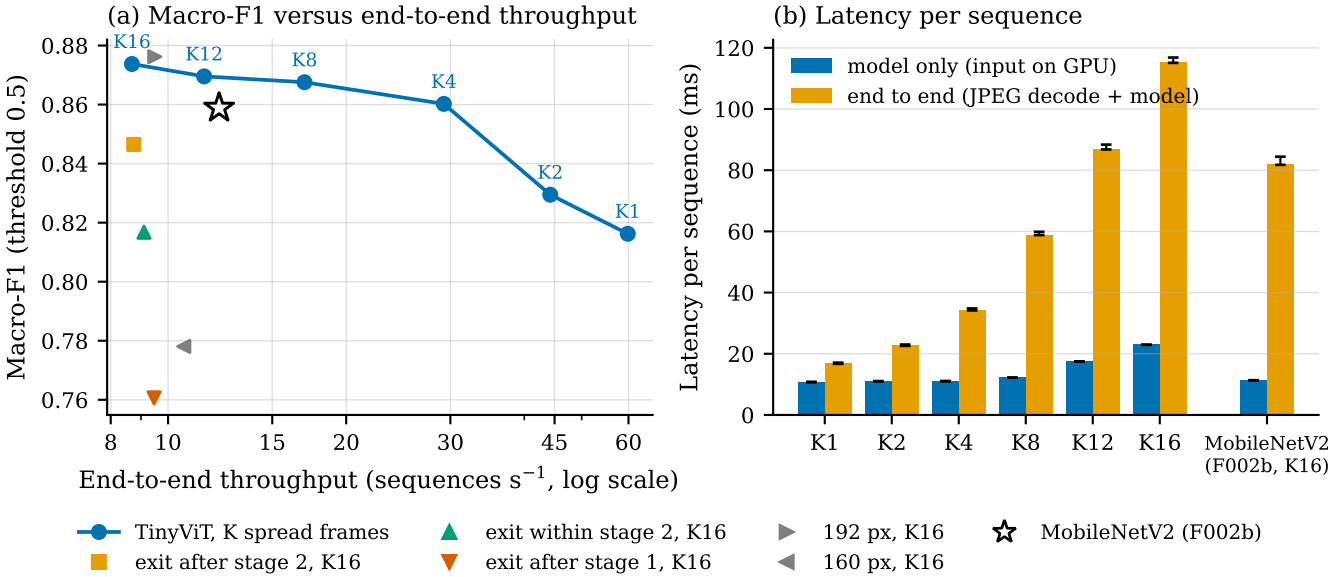
Runtime on a GTX 1080 Ti
(a) Macro-F1 versus end-to-end throughput
Macro-F1 (threshold 0.5)
0.88
0.86
0.84
0.82
0.80
0.78
0.76
K16
K12
K8
K4
K2
K1
8
10
15
20
30
45
60
End-to-end throughput (sequences s^{-1}, log scale)
(b) Latency per sequence
Latency per sequence (ms)
120
100
80
60
40
20
0
model only (input on GPU)
end to end (JPEG decode + model)
K1
K2
K4
K8
K12
K16
MobileNetV2
(F002b, K16)
TinyViT, K spread frames
exit after stage 2, K16
exit within stage 2, K16
exit after stage 1, K16
192 px, K16
160 px, K16
MobileNetV2 (F002b)