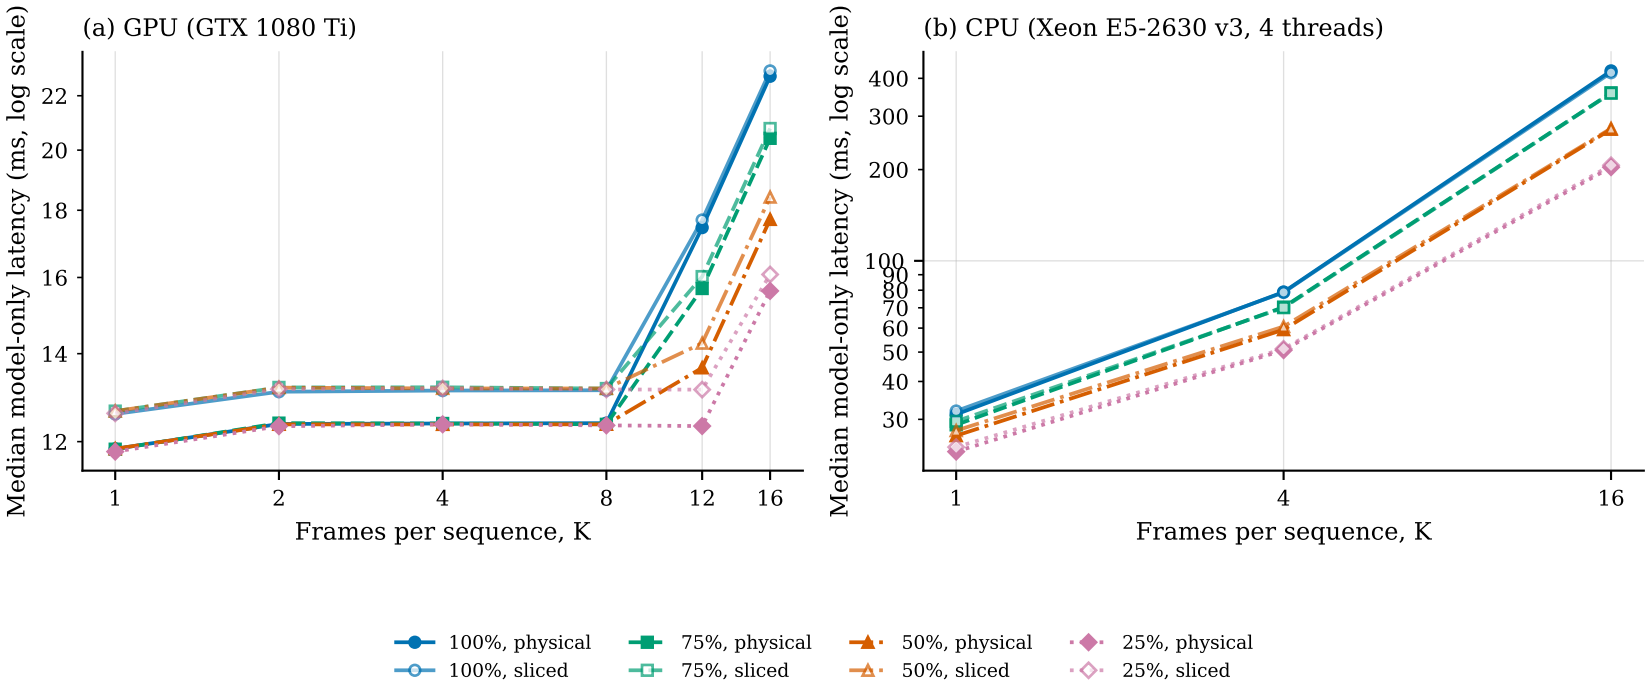
Latency of the nested pruned widths
(a) GPU (GTX 1080 Ti)
(b) CPU (Xeon E5-2630 v3, 4 threads)
Median model-only latency (ms, log scale)
Frames per sequence, K
100%, physical
100%, sliced
75%, physical
75%, sliced
50%, physical
50%, sliced
25%, physical
25%, sliced